



Deep Learning Brain Tumor Classification

Group 1: Medical Image Classifiers
Anulitha Vardhan
Carsten Lingemann
Sundari Kala

Problem Description

- We are leveraging 3 different articles that each used a different pre-trained model (VGG16, ResNet50, InceptionV3) to classify Brain Tumor Images.
- We are also aiming to create a general medical image classification model that can be used to diagnose other medical conditions, such as leukemia, as well as chest x-rays detecting pneumonia and COVID-19.

Motivation

- Classification of medical images is important because it can lead to a correct or incorrect diagnosis and treatment plan.
- An untreated condition can continue to grow and cause irreversible damage.
- Medical images analysis can be a very useful research idea and can be a great career opportunity in the future.
- Building deep learning models for medical image analysis involves a diverse range of skills and knowledge, including computer science, mathematics, statistics, and biology.

Existing Approaches

Inception-V3

- Comparing accuracy of the Inception V3 performance with the VGG-16.
- 100 images in each category (normal/abnormal) are selected from the database of 130 subjects.
- Average accuracy of inception architecture is 95.1%, whereas VGG-16 is 92.8%
- [Link to article](#)

ResNet-50

- Used [ResNet](#)-50 to extract features of pre-preprocessed (normalized and resized) images, and fed them into a feature selection algorithm.
- Selected features used to train a SVM classifier. Learning rate, weight decay, and optimizer to improve performance.
- Used Grad-CAM to visualize the regions of the image that contributed the most towards the model's prediction.

VGG-16


- Used 23-layer CNN for multiclass brain tumor classification (meningioma, glioma, pituitary) of 3064 images with 97.8% accuracy
- 23-layer CNN did not perform well for binary brain tumor classification (normal, abnormal) of only 152 images, so they used pretrained VGG-16 model with 100% accuracy
- [Link to article](#)

3 Different Data Sets


Data Set	Classes	Training	Validation	Test	Total
Brain Tumor	<ul style="list-style-type: none">• Glioma Tumor• Meningioma Tumor• Pituitary Tumor• Normal	4570	1142	1311	7023
Leukemia	<ul style="list-style-type: none">• ALL - Acute Lymphoblastic Leukemia• HEM - Normal	6090	1260	1263	8613
Chest X-Rays	<ul style="list-style-type: none">• COVID-19• Pneumonia-Bacterial• Pneumonia-Viral• Normal	4603	1841	2764	9208

Example Images

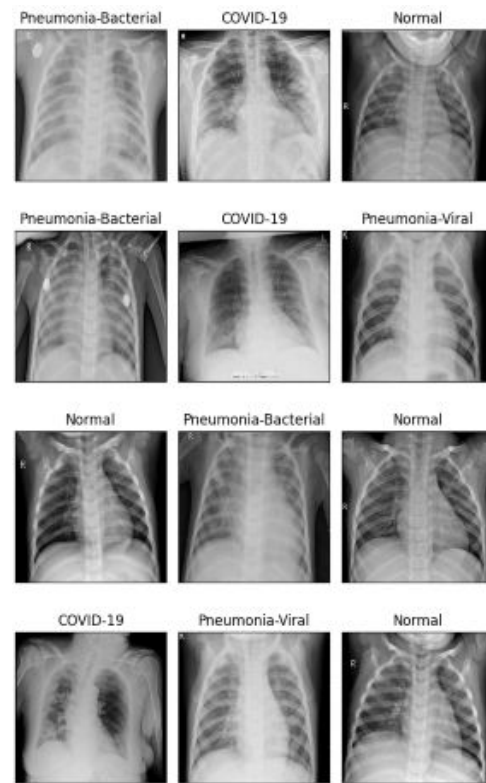
Brain Tumor



Leukemia



Chest X-Ray




VGG-16 Image Augmentation

Transformations Applied:

- RandomZoom(0.3)
- RandomFlip(“horizontal”)
- RandomWidth(0.2)
- RandomHeight(0.2)
- RandomFlip(“vertical”)

Leukemia Augmentation Example



VGG-16 Experiments / Observations

Model: "bt_baseline"






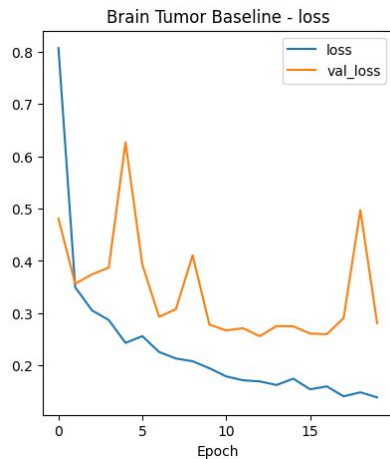
Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 224, 224, 3)]	0
sequential_1 (Sequential)	(None, None, None, 3)	0
resizing_1 (Resizing)	(None, 224, 224, 3)	0
tf.__operators__.getitem_1 (SlicingOpLambda)	(None, 224, 224, 3)	0
tf.nn.bias_add_1 (TFOpLambda)	(None, 224, 224, 3)	0
vgg16 (Functional)	(None, 7, 7, 512)	14714688
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 512)	0
dense_4 (Dense)	(None, 1024)	525312
dense_5 (Dense)	(None, 1024)	1049600
dense_6 (Dense)	(None, 512)	524800
dropout_1 (Dropout)	(None, 512)	0
dense_7 (Dense)	(None, 4)	2052

=====
Total params: 16,816,452
Trainable params: 2,101,764
Non-trainable params: 14,714,688


- Adding Batch Normalization layer between each dense function improved the Chest X-Ray test accuracy by $\sim 1\%$, but decreased the test accuracy on the Brain Tumor and Leukemia data sets
- Using Binary Cross-entropy for the Leukemia data set, along with the sigmoid activation function on the output layer, and RMSProp optimizer improved the test accuracy by $\sim 7\%$
- Increasing epochs from 25 to 50 led to overfitting, and reduced the test accuracy on all 3 data sets

VGG-16 Results

Test Accuracy	
Brain Tumor	91.99%
Leukemia	69.20%
Chest X-Ray	84.19%



VGG-16 Predictions




Inception-V3 Image Augmentation

Transformations Applied:

- RandomZoom(0.2)
- RandomFlip("horizontal")
- RandomWidth(0.1)
- RandomHeight(0.1)
- RandomFlip("vertical")
- RandomCrop(height=180, width=180)

Chest X-Ray Augmentation Example



Inception-V3 Experiments / Observations

- Tried different optimizers, activation and loss functions.
- Added the Batch Normalization to improve the performance, however it didn't predicted the values accurately.
- The following activation, loss and optimizers performed well for my model.
 - Activation - softmax
 - Optimizer - Adam
 - Loss function - sparse categorical crossentropy

Model: "model_1"


Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 224, 224, 3)]	0
sequential_2 (Sequential)	(None, 180, 180, 3)	0
resizing_1 (Resizing)	(None, 224, 224, 3)	0
tf.math.truediv_1 (TFOpLambda)	(None, 224, 224, 3)	0
tf.math.subtract_1 (TFOpLambda)	(None, 224, 224, 3)	0
inception_v3 (Functional)	(None, 5, 5, 2048)	21802784
max_pooling2d_5 (MaxPooling2D)	(None, 2, 2, 2048)	0
flatten_1 (Flatten)	(None, 8192)	0
dense_2 (Dense)	(None, 512)	4194816
dropout_1 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 4)	2052
Total params: 25,999,652		
Trainable params: 4,196,868		
Non-trainable params: 21,802,784		

Inception-V3 Results



Test Accuracy

Brain Tumor	86.12%
Leukemia	75.14%
Chest X-Ray	75.58%



Brain Tumor



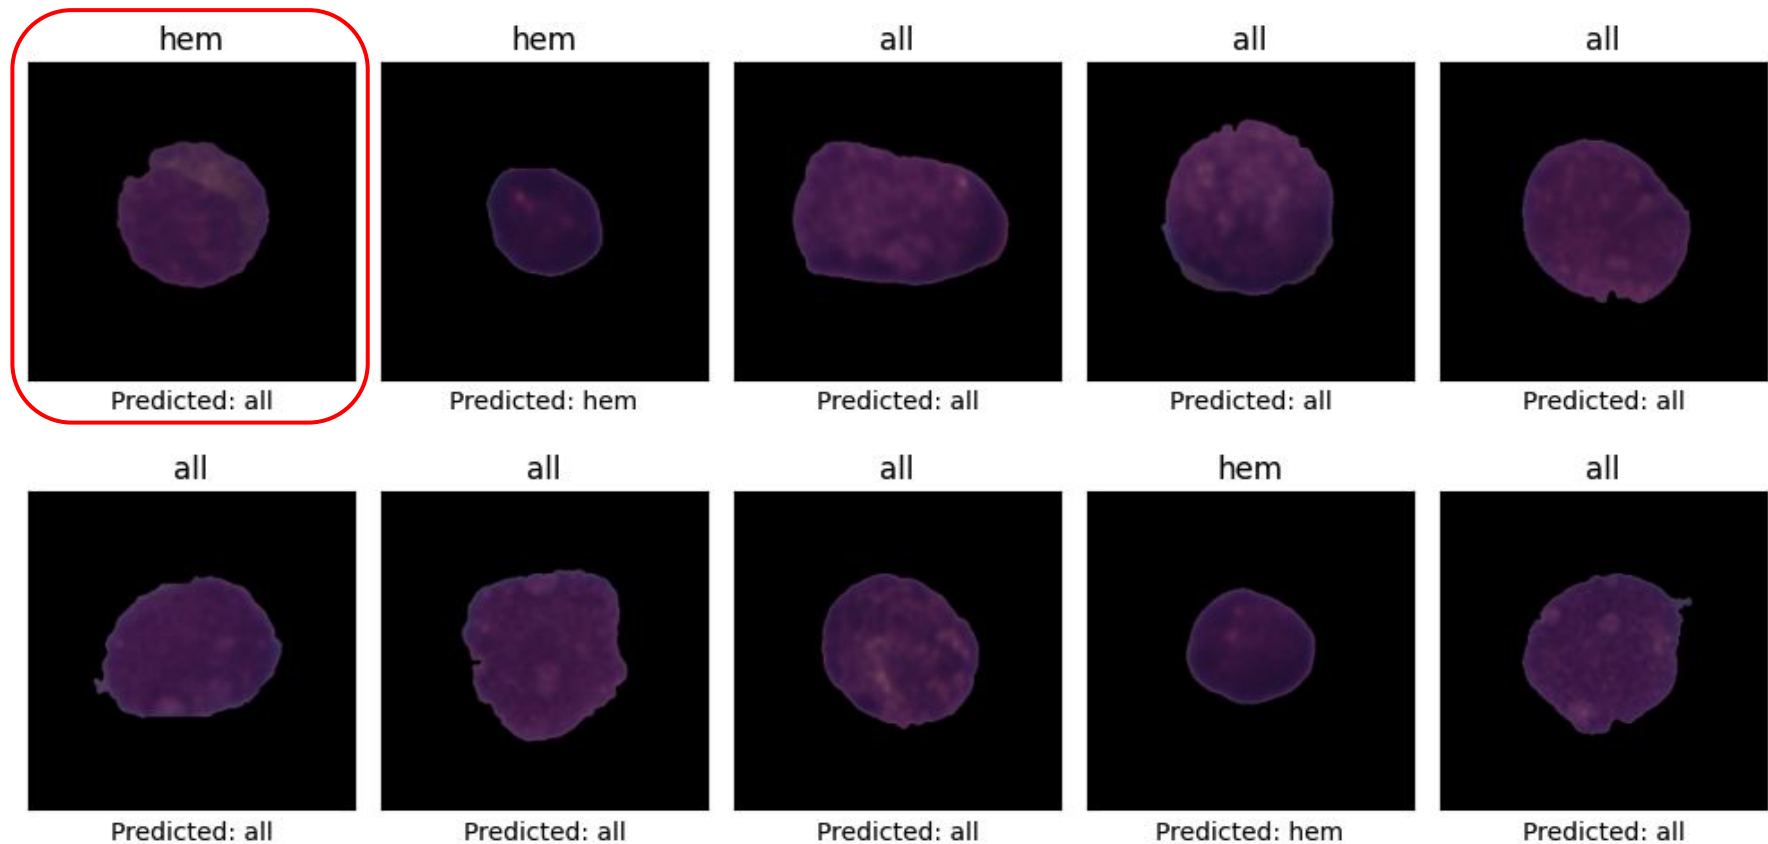
Leukemia



Chest X-Ray



Inception-V3 Predictions




ResNet-50 Image Augmentation

Transformations Applied:

- rotation_range=10,
- width_shift_range=0.05,
- height_shift_range=0.05,
- horizontal_flip=True

Leukemia Augmentation Example



ResNet-50 Experiments / Observations

```
global_average_pooling2d (GlobalAveragePooling2D) 0 ['conv5_block3_out[0][0]']
dropout (Dropout) (None, 2048) 0 ['global_average_pooling2d[0][0]']
dense (Dense) (None, 4) 8196 ['dropout[0][0]']
=====
Total params: 23,595,908
Trainable params: 23,542,788
Non-trainable params: 53,120
```

Complicated architecture reduced val-accuracy

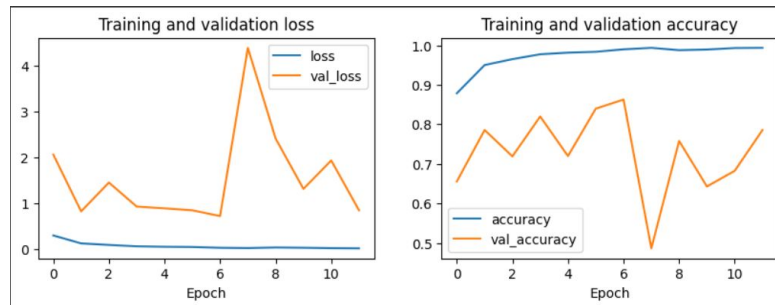
- Training the entire Resnet [unfrozen] with early stopping, and LROnPlateau with only Global Average pooling layer, Dropout and Dense layers yielded 97% val_accuracy and 98.09% test accuracy.
- The same Resnet50 model performed decently well on XRay and Leukemia datasets with 84 and 86% accuracy respectively.
- MRI images performed poorly with complex architecture on Leukemia and Chest X-Ray datasets and excelled with a simpler architecture with ES and LRoP.

ResNet-50 Results


Test Accuracy

Brain Tumor	98.93%
Leukemia	86.27%
Chest X-Ray	84.55%


Leukemia



Brain Tumor




Chest X-ray




ResNet-50 Predictions


True label: notumor
Predicted label: notumor




True label: meningioma
Predicted label: meningioma




True label: notumor
Predicted label: notumor




True label: meningioma
Predicted label: meningioma




True label: glioma
Predicted label: glioma




True label: pituitary
Predicted label: pituitary




True label: meningioma
Predicted label: meningioma




True label: meningioma
Predicted label: meningioma



True label: notumor
Predicted label: notumor



True label: pituitary
Predicted label: pituitary



Challenges

- Existing research papers were using methods that are now deprecated (i.e. ImageDataGenerator), so it took some time to replicate their results using the newer methods (image_dataset_from_directory) and not all of the prior features were available in the newer methods
- Maxing out GPU consumption on GoogleColab
- Ran out of time to experiment with Albumentations library for image augmentation
- Figuring out which optimizer to use for each pre-trained model and data set
- ResNet50 performed poorly on complex architectures, & low information available on Brain MRI classification with 4 categories

Future Work

- Create a custom loss function to more heavily penalize false negatives, since the consequence of missing a condition is much worse than getting a false positive
- Experiment with Albumentations image augmentation library
- Customize the image augmentation pre-processing for each dataset (i.e. crop and apply brightness specifically for the leukemia data set)
- Experiment with combining all 3 pre-trained models
- Experiment with using un-frozen layers on other pre-trained models
- Could potentially leverage more advanced techniques (i.e. transformers)

Conclusion

- Could not generalize a model to work with all 3 data sets:
 - ResNet-50 performed the best on all 3 models overall
 - Inception-V3 performed the best on the Leukemia data set
 - VGG-16 performed the best on the Chest X-Ray data set
- Different pre-trained models could work better on different types of data sets
- Different architectures (including loss functions, optimizers, and activations) are necessary for binary classification vs. multi classification

References

- Nickparvar, Msoud. (2021, September 24). *Brain Tumor MRI Dataset* [Data set]. Kaggle. <https://www.kaggle.com/datasets/masoudnickparvar/brain-tumor-mri-dataset>
- Khan, Md Saikat Islam et al. "Accurate brain tumor detection using deep convolutional neural network." *Computational and structural biotechnology journal* vol. 20 4733-4745. 27 Aug. 2022, doi:10.1016/j.csbj.2022.08.039
- Sadad, Tariq, et al. "Brain tumor detection and multi-classification using advanced deep learning techniques." *Microscopy Research and Technique* 84.6 (2021): 1296-1308.
- Baloni, Dev, and Shashi Kant Verma. "Detection of hydrocephalus using deep convolutional neural network in medical science." *Multimedia Tools and Applications* 81.12 (2022): 16171-16193.
- Elmorshedy, Abeer. (2022, December 10). *leukemia* [Data set]. Kaggle. <https://www.kaggle.com/datasets/abeerelmorshedy/leukemia>
- R Tamilarasi and S Gopinathan 2021 J. Phys.: Conf. Ser. 1964 072022